

Image Classification



64x64

→ Cat? (0/1)

Object detection



Neural Style Transfer




Image Source: deeplearning.ai

Computer Vision Problems

Classical Computer Vision Pipeline

CV experts

1. Select / develop features: SURF, HoG, SIFT, RIFT, ...
2. Add on top of this Machine Learning for multi-class recognition and train classifier




Classical CV feature definition is domain-specific and time-consuming

Neural Network



Warren
McCulloch



Walter Pitts

A LOGICAL CALCULUS OF THE IDEAS IMMANENT IN NERVOUS ACTIVITY,
Bulletin of Mathematical Biophysics, Vol. 5, pp. 115-133 (1943).

Neural Network

Here x_1 and x_2 are normalized attribute value of data.

y is the output of the neuron , i.e the class label.

x_1 and x_2 values multiplied by weight values w_1 and w_2 are input to the neuron x .

Value of x_1 is multiplied by a weight w_1 and values of x_2 is multiplied by a weight w_2 .

Given that

- $w_1 = 0.5$ and $w_2 = 0.5$
- Say value of x_1 is 0.3 and value of x_2 is 0.8,
- So, weighted sum is :
- $\text{sum} = w_1 \times x_1 + w_2 \times x_2 = 0.5 \times 0.3 + 0.5 \times 0.8 = 0.55$





Fig1: an artificial neuron


Why We Need Multi Layer ?

Linear Separable:



Linear inseparable:

$x \nabla y$



Edge Detection



Vertical edges




Horizontal edges

How do we detect these edges


Neural Network?

- ❑ Suppose an image is of the size $68 \times 68 \times 3$
 - Input feature dimension then becomes 12,288
- ❑ If Image size is of $720 \times 720 \times 3$
 - Input feature dimension becomes 1,555,200
- ❑ Number of parameters will swell up to a HUGE number
- ❑ Result in more computational and memory requirements



Another Application

Digit Recognition



$X_1, \dots, X_n \in \{0,1\}$ (Black vs. White pixels)

$Y \in \{5,6\}$ (predict whether a digit is a 5 or a 6)

The Bayes Classifier

In class, we saw that a good strategy is to predict:

$$\arg \max_Y P(Y|X_1, \dots, X_n)$$

- (for example: what is the probability that the image represents a 5 given its pixels?)

So ... how do we compute that?

The Bayes Classifier

Use Bayes Rule!

$$P(Y|X_1, \dots, X_n) = \frac{P(X_1, \dots, X_n|Y)P(Y)}{P(X_1, \dots, X_n)}$$

Likelihood Prior
Normalization Constant

Why did this help? Well, we think that we might be able to specify how features are “generated” by the class label

The Bayes Classifier

Let's expand this for our digit recognition task:

$$P(Y = 5|X_1, \dots, X_n) = \frac{P(X_1, \dots, X_n|Y = 5)P(Y = 5)}{P(X_1, \dots, X_n|Y = 5)P(Y = 5) + P(X_1, \dots, X_n|Y = 6)P(Y = 6)}$$
$$P(Y = 6|X_1, \dots, X_n) = \frac{P(X_1, \dots, X_n|Y = 6)P(Y = 6)}{P(X_1, \dots, X_n|Y = 5)P(Y = 5) + P(X_1, \dots, X_n|Y = 6)P(Y = 6)}$$

To classify, we'll simply compute these two probabilities and predict based on which one is greater

Model Parameters

For the Bayes classifier, we need to “learn” two functions, the likelihood and the prior


How many parameters are required to specify the prior for our digit recognition example?



Model Parameters

How many parameters are required to specify the likelihood?

- (Supposing that each image is 30x30 pixels)



Drive into CNN


In a convolutional network (ConvNet), there are basically three types of layers:

1. Convolution layer
2. Pooling layer
3. Fully connected layer



A convolutional layer

A CNN is a neural network with some convolutional layers (and some other layers). A convolutional layer has a number of filters that does convolutional operation.



Convolution

These are the network parameters to be learned.

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

6 x 6 image

1	-1	-1
-1	1	-1
-1	-1	1

Filter 1

-1	1	-1
-1	1	-1
-1	1	-1

Filter 2

⋮ ⋮

Each filter detects a small pattern (3 x 3).

Convolution


1	-1	-1
-1	1	-1
-1	-1	1

Filter 1

stride=1

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

Dot
product



6 x 6 image

Convolution


1	-1	-1
-1	1	-1
-1	-1	1

Filter 1

If stride=2

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

6 x 6 image



Convolution

1	-1	-1
-1	1	-1
-1	-1	1

Filter 1

stride=1

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

6 x 6 image

3	-1	-3	-1
-3	1	0	-3
-3	-3	0	1
3	-2	-2	-1

Convolution

-1	1	-1
-1	1	-1
-1	1	-1


Filter 2

stride=1

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

6 x 6 image


Repeat this for each filter



Two 4 x 4 images
Forming 4 x 4 x 2
matrix

Convolution over Volume

Color image



1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0


1	-1	-1
-1	1	-1
-1	-1	1

Filter 1


-1	1	-1
-1	1	-1
-1	1	-1


Filter 2

Convolution v.s. Fully Connected





Fully-
connected






fewer parameters!





Fewer parameters

Even fewer parameters




Suppose we have 10 filters applying on input (6 X 6 X 3), each of shape 3 X 3 X 3. What will be the number of parameters in that layer?

- Number of parameters for each filter = $3*3*3 = 27$
- There will be a bias term for each filter, so total parameters per filter = 28
- As there are 10 filters, the total parameters for that layer = $28*10 = 280$



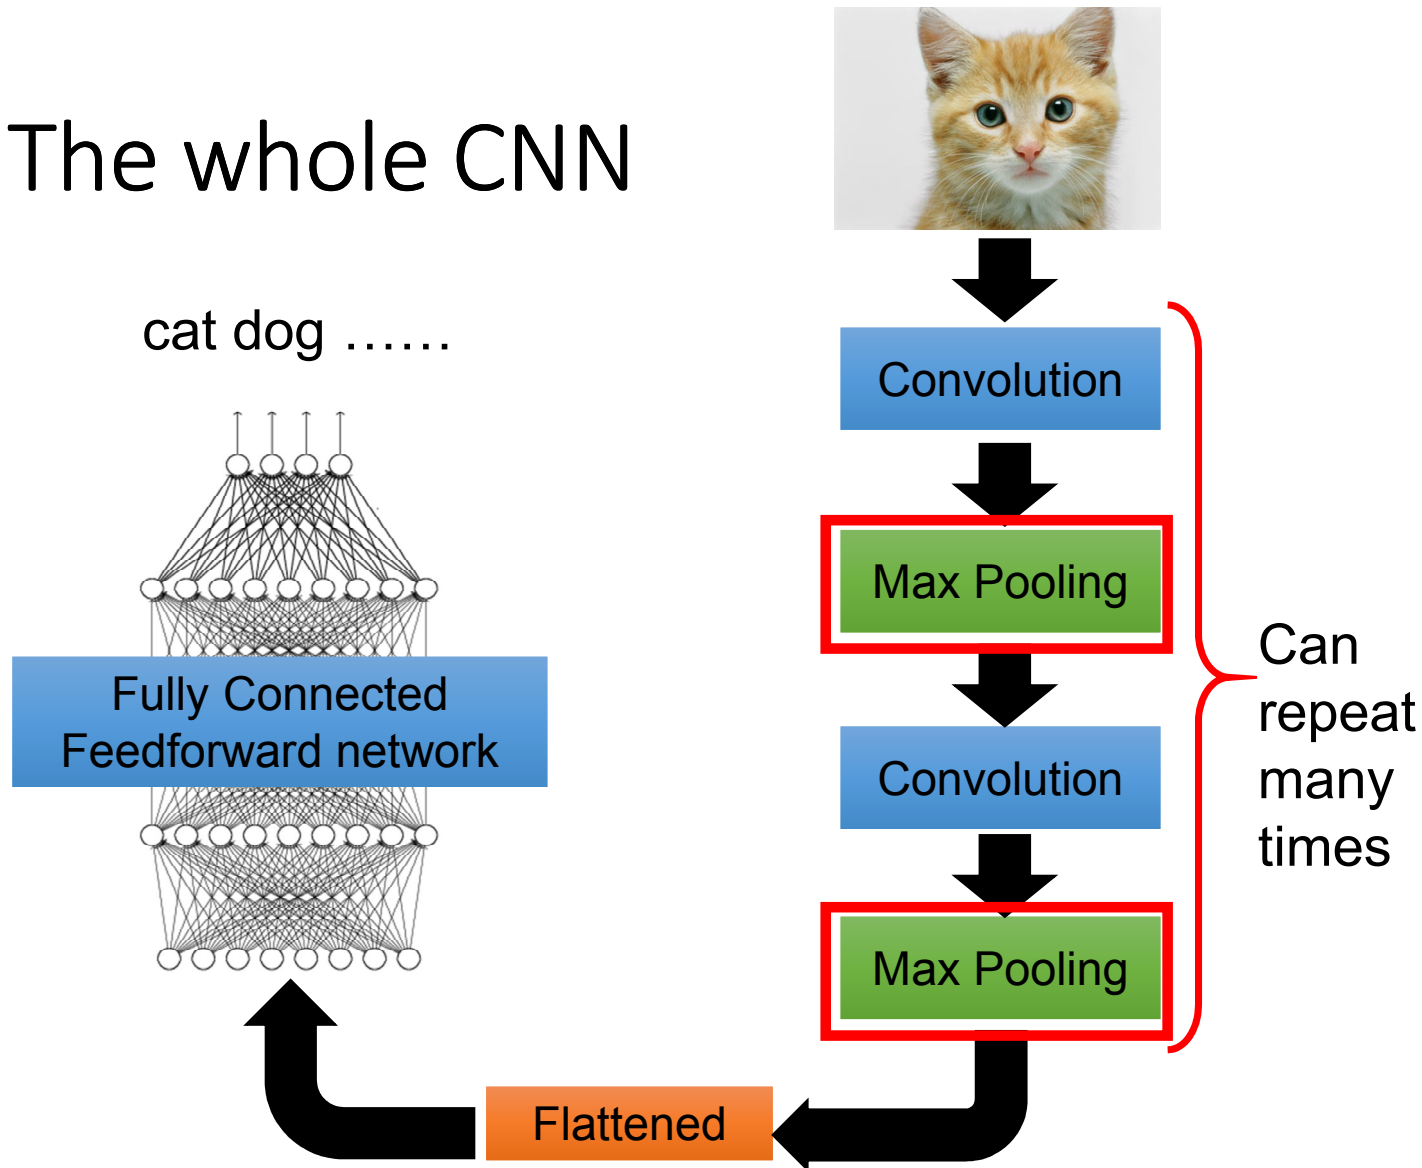
QUIZ

Simple Convolutional Neural Network



- Size of feature vector : $(n+2p-f)/s + 1$
- n : dimension of matrix
- p : size of padding
- f : size of filter
- s : size of stride

The whole CNN



Max Pooling

1	-1	-1
-1	1	-1
-1	-1	1

Filter 1

-1	1	-1
-1	1	-1
-1	1	-1

Filter 2

3	-1	-3	-1
-3	1	0	-3
-3	-3	0	1
3	-2	-2	-1

-1	-1	-1	-1
-1	-1	-2	1
-1	-1	-2	1
-1	0	-4	3

Why Pooling

- Subsampling pixels will not change the object

bird



Subsampling


bird



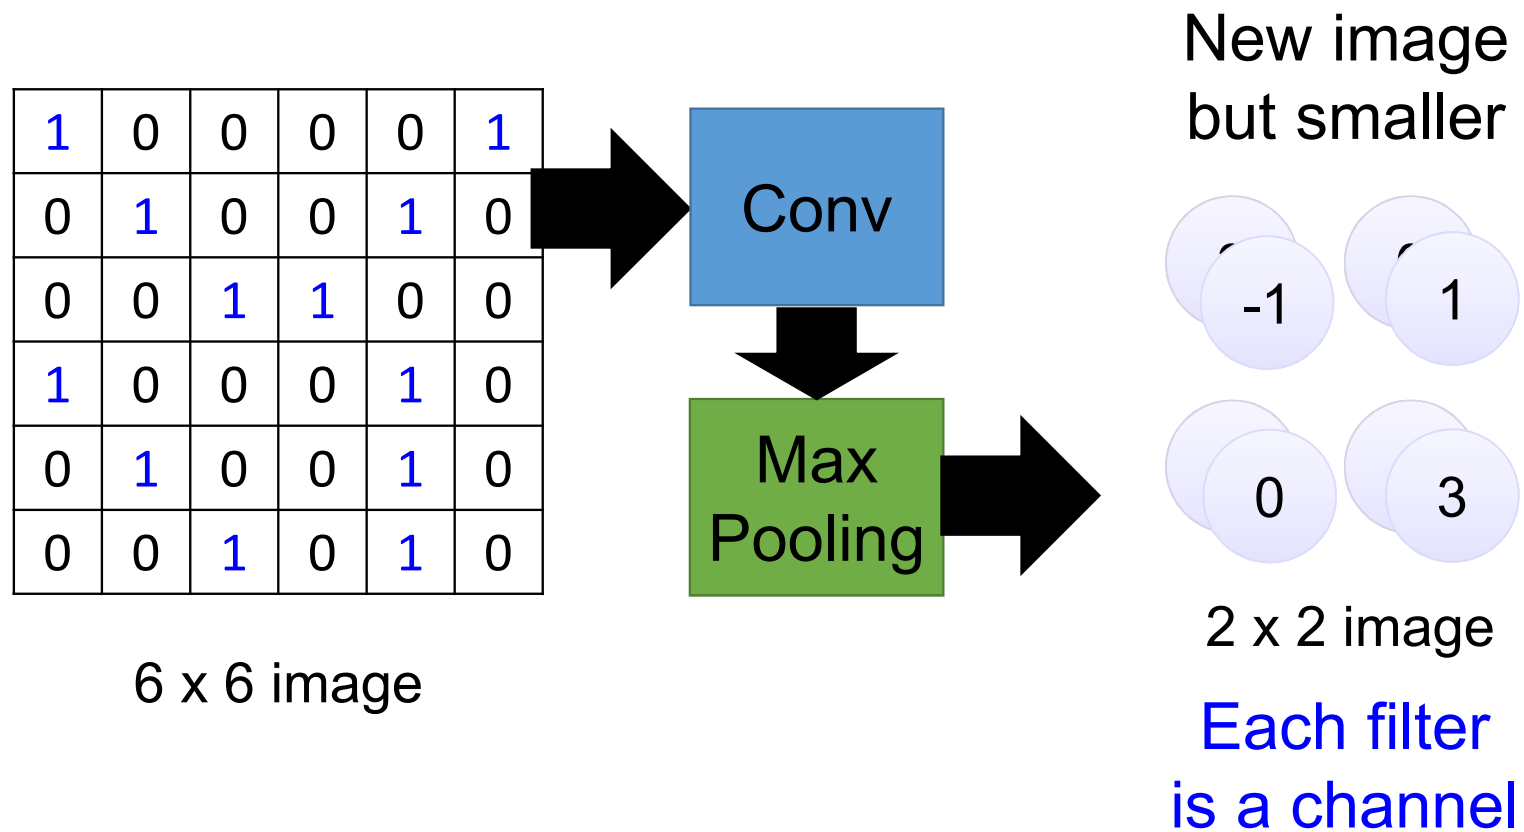
We can subsample the pixels to make image
smaller → fewer parameters to characterize the image

A CNN compresses a fully connected network in two ways:


- Reducing number of connections
- Shared weights on the edges
- Max pooling further reduces the complexity



Max Pooling



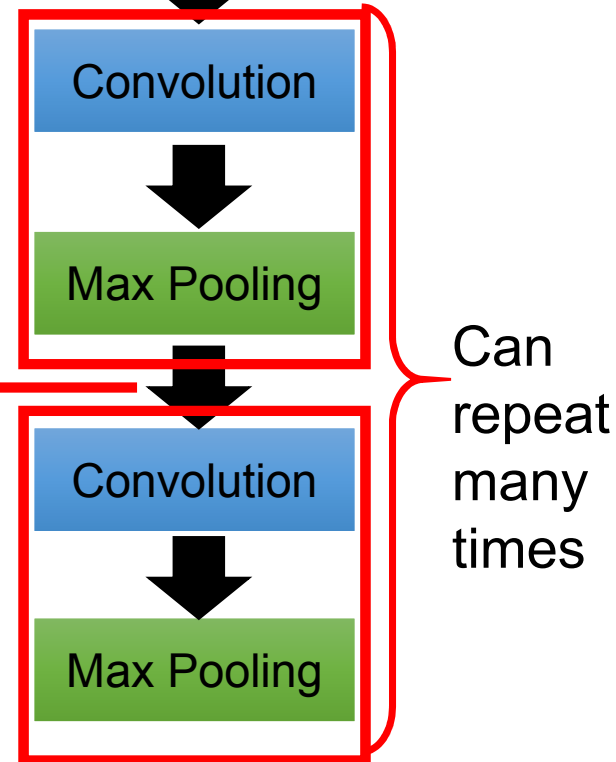

The whole CNN



A new image

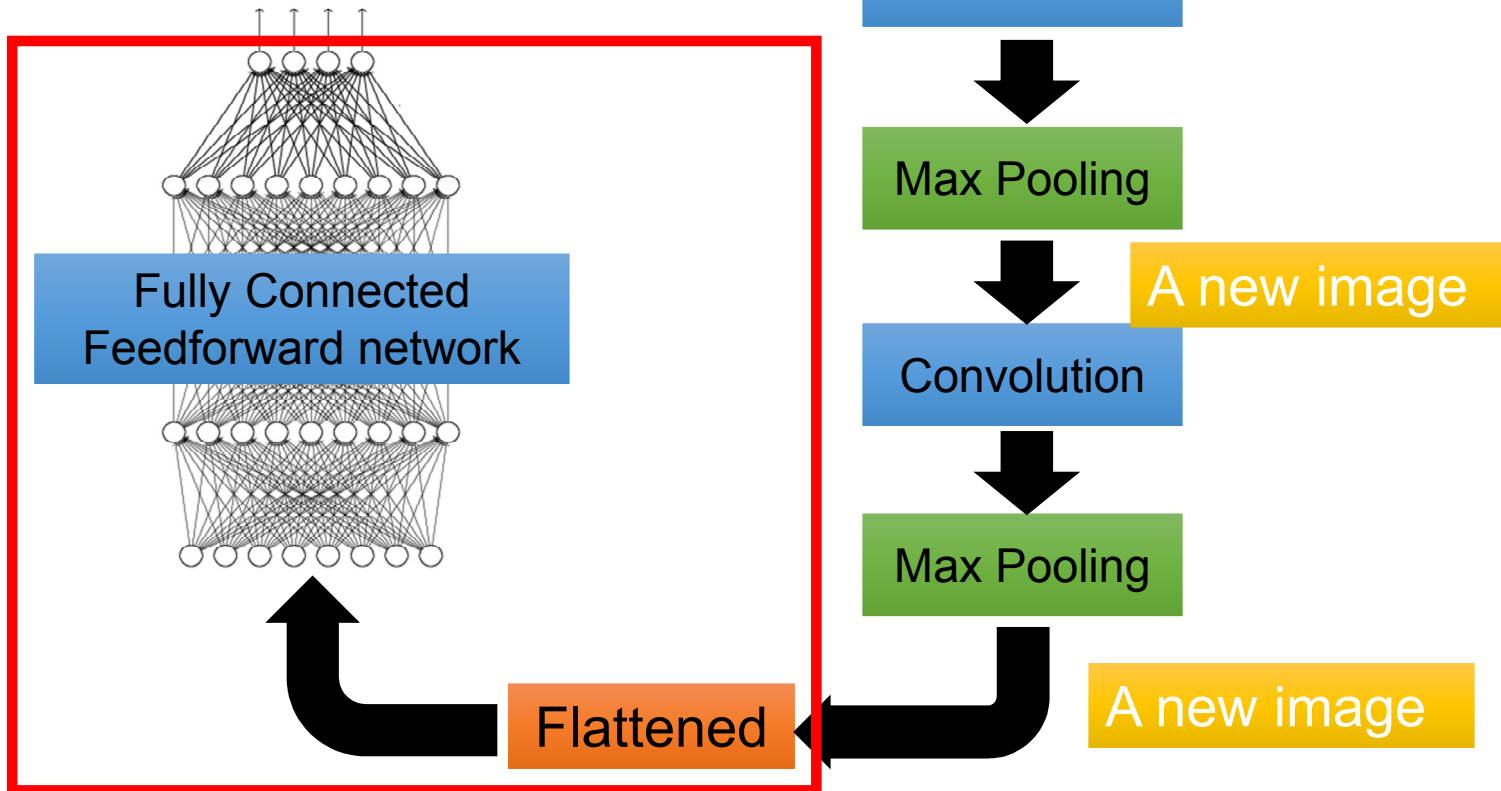
Smaller than the original image

The number of channels is the number of filters

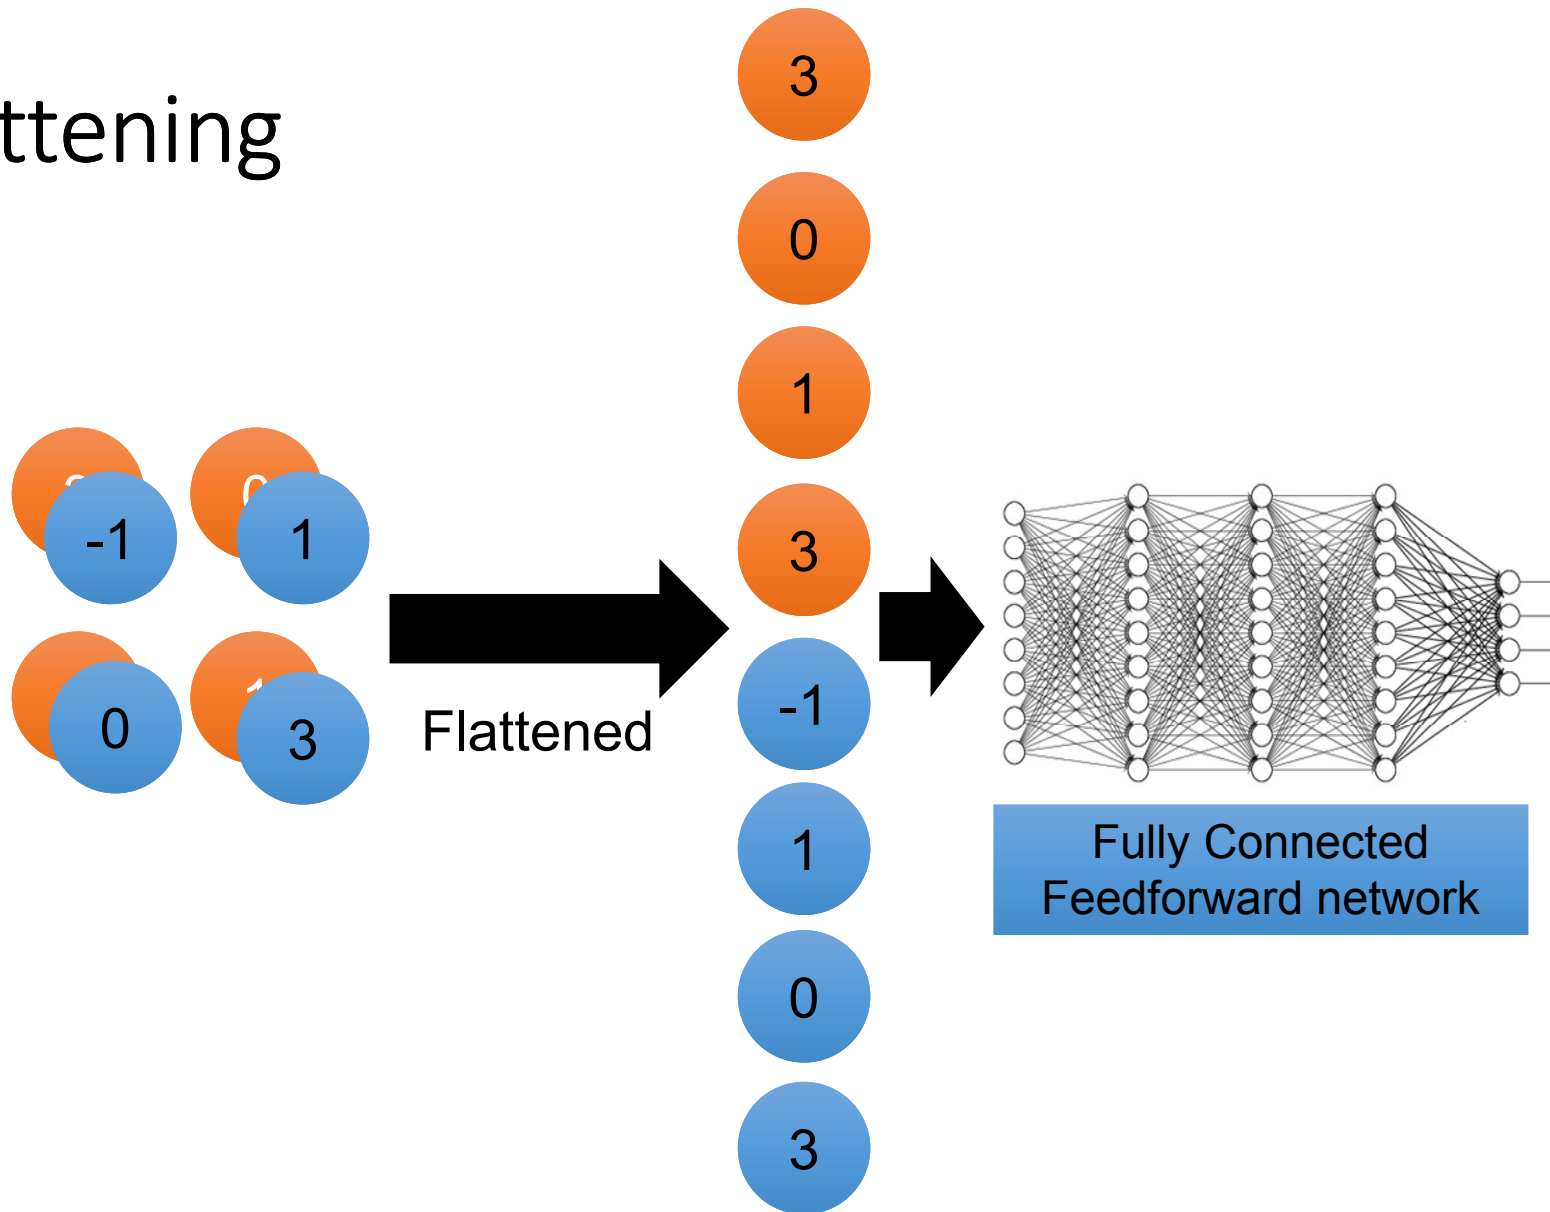


The whole CNN

cat dog




Flattening



Classic Networks


1. LeNet-5
2. AlexNet
3. VGG

LeNet-5



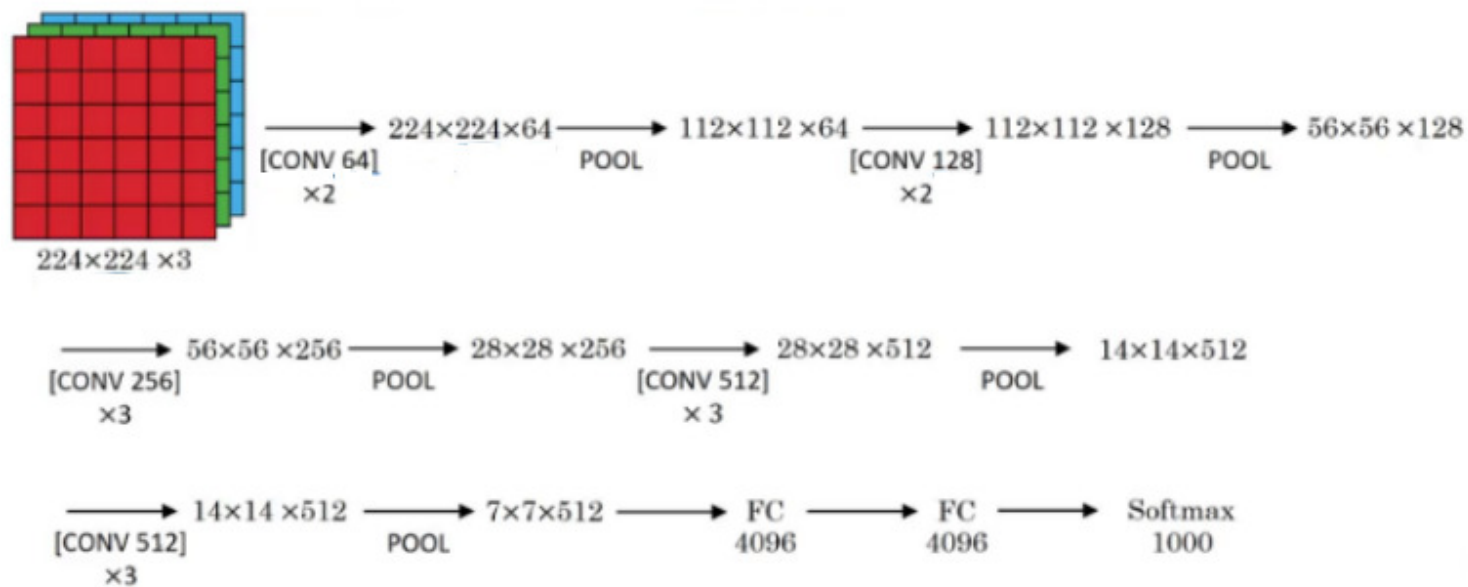
- **Parameters:** 60k
- **Layers flow:** Conv \rightarrow Pool \rightarrow Conv \rightarrow Pool \rightarrow FC \rightarrow FC \rightarrow Output
- **Activation functions:** Sigmoid/tanh and ReLU

AlexNet



- Parameters:** 60 million
- Activation functions:** ReLU

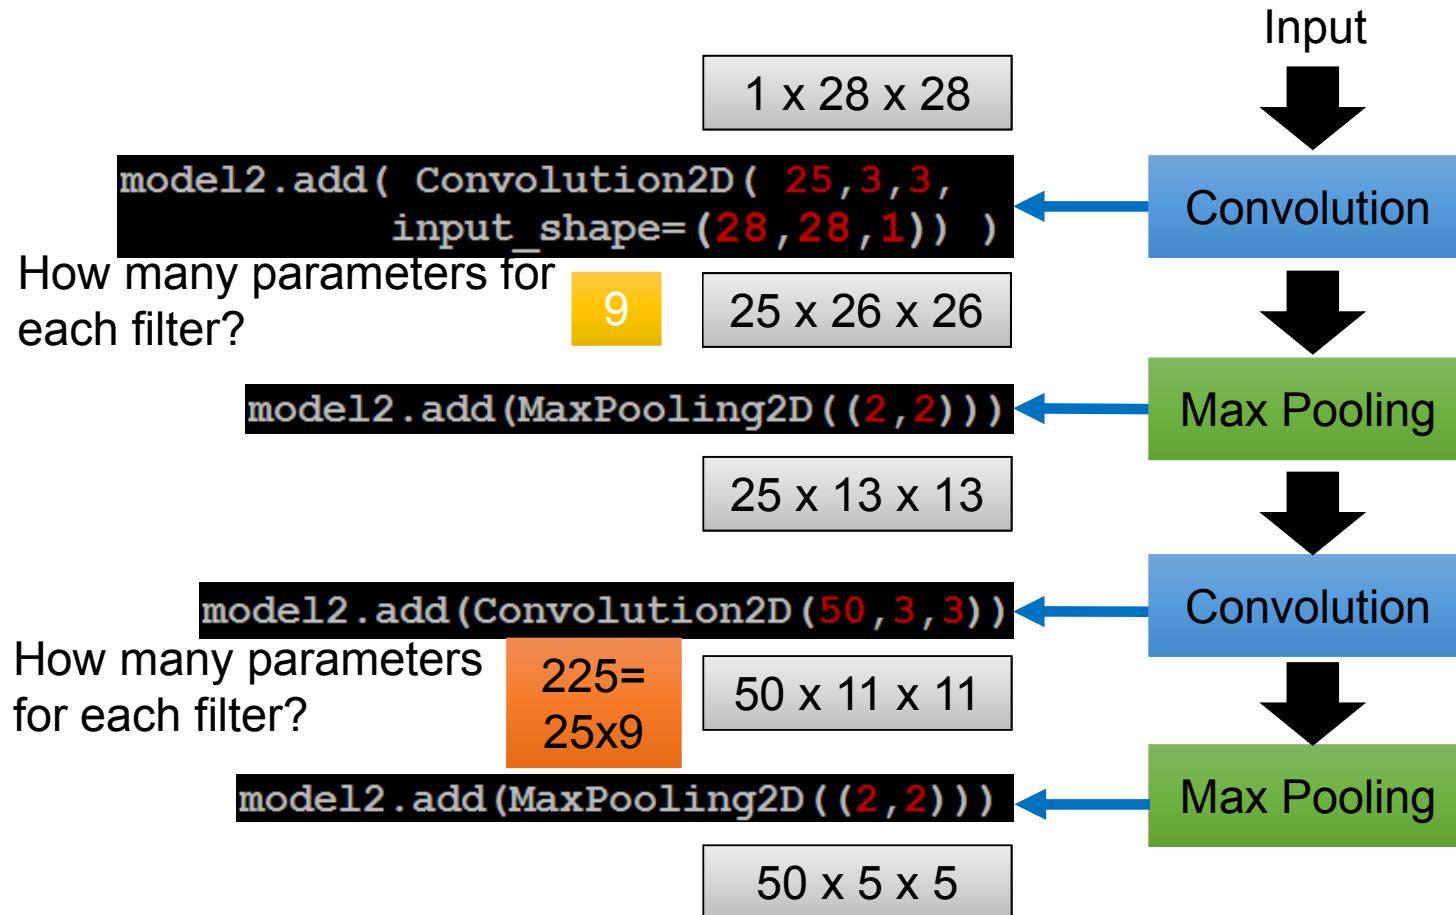
VGG-16



- Parameters: 138 million
- Pool: MAX with stride 2
- CONV layer: stride 1

CNN in Keras

Only modified the *network structure* and *input format (vector -> 3-D array)*



CNN in Keras

Only modified the *network structure* and *input format (vector -> 3-D array)*

